



Linux Systems and Open Source Software

Basics of Performance Analysis Part I

Chia-Heng Tu

Dept. of Computer Science and Information
Engineering

National Cheng Kung University

Fall 2022

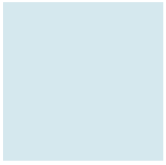




Outline

- What is Performance?
- Time Measurement for Performance Analysis
- Standard Time Measurement (POSIX)
- Native Linux Time Measurement
- Time Measurement for x86 Platforms





WHAT IS PERFORMANCE?





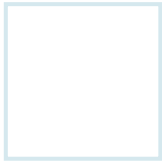
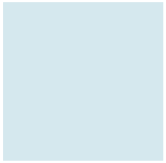
Define *Performance*

Plane	<i>Speed</i>	<i>DC to Paris time</i>	<i>Passengers</i>	<i>Throughput (p x mph)</i>
Boeing 747	610 mph	6.5 hours	470	286,700
Concorde	1350 mph	3 hours	132	178,200

Choose an airplane with better *performance*

- To pick a fastest airplane?
- To pick a high-throughput airplane?



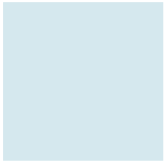


Define *Performance* (Cont'd)

Plane	Speed	DC to Paris time	Passengers	Throughput (p x mph)
Boeing 747	610 mph	6.5 hours	470	286,700
Concorde	1350 mph	3 hours	132	178,200

- **Time** to do the task
 - **Execution time**, response time, latency
- **Tasks** per day, hour, week, sec, ns. ..
 - **Throughput**, bandwidth






Performance Measurement

- In CS, when talk about **performance**, we should note...
- What is the **goal** here?
 - A well-tweaked program?
 - A fast machine?
 - ?
- **How** is performance measured?
 - Manually or automatically; will discussed in this course
- Example:
How to **choose** amongst different machines?
 - Cost: price, technology metrics
 - Performance
 - **Metrics**: **time** and **processing speed**; indicate relative performance
 - E.g., **run time** or **X's per second**



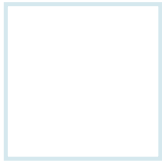
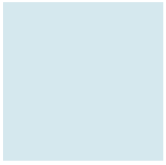


Performance Metrics



- What is a **metric**?
 - Basis for measuring
 - Basis for comparison
- Metric varies across different situations
 - Job: Salary, Responsibilities; School: Grades;
 - Mutual Funds: Total Return, Risk
 - Cars: Top speed, Acceleration, Impact test
- Metrics are used to **Order** and **Compare**
- Many metrics are possible, many used





Computer Performance Metrics

- **Execution Time**

- CPU time, wall clock time

- Millions of Instructions Per Second (MIPS)

- Machine instructions
small primitive units generated by the compiler

- Cycles Per Instruction (CPI)

- Fixed length time periods, normalization for technology

- Clock Rate (Megahertz)

- Millions of cycles per second

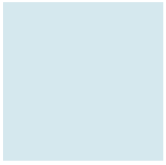
=> Many metrics are used

- Some better, some worse

- **Goal:** use metrics which **reflect performance** delivered by the underlying machine to real user programs, real applications

- Quantize the **combined effect** of the application (SW) and machine (HW)

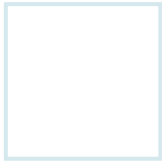
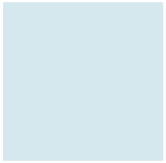




Measuring Execution Time

- Steps
 1. Select a program
 2. Execute the program
 3. Measure the CPU time, or the wall clock (stopwatch) time
- Performance = $1 / (\text{Execution time})$
 - Smaller execution time -> better performance
 - Larger execution time -> worse performance
- Questions?
 - What to be compared? (the performance of different programs?)
 - How to measure time?





Programs under Test (What)

- Target program
 - Would like to run EXACTLY the **same program**, but...
 - Programs are large and unwieldy, input data is critical, and they and their usage would change over time
 - Need to **run entire application programs at full data set sizes** to get good performance information
 - Best **predictors**, tricky, difficult, sometimes misleading
- “*Benchmarking*”
 - How applications will really run
 - Lies? and Benchmarks
 - Common benchmark programs are used to evaluate the performance amongst **commercial machines**
 - [SPEC](#), TPC, Dhrystones, [Antutu](#), etc.





Time Measurement (How)

- How to measure execution time?
 - Watch (start/stop the meter; wall clock)
 - Computer timing: User (processor time) + System (operating system -- I/O, etc.)
 - Compare based on *processor time* for *processor performance*
 - Important, but of decreasing importance for SYSTEM performance
 - Concurrent users? Measurement errors?
(Involves **What** to measure!!!)
- Usually, you should know **What** and then you will know **How**





Time

- **CPU Execution Time**
 $\text{= CPU clock cycles} * \text{Clock cycle time}$
 $\text{= CPU clock cycles} / \text{Clock rate}$
- **Elapsed time (or wall-clock time)**
 - It is the actual time taken between the start and the end of the task
- **CPU Execution Time does not necessarily mean the wall-clock time**
 - I.e., the elapsed time for running the program may be different from the time that CPU spent on the program
- Every conventional processor has a clock with an associated clock cycle time or clock rate
- Every program runs in an integer number of clock cycles
 $x \text{ MHz} = x \text{ millions of cycles/second (clock rate)}$
 $1 / (x \text{ MHz}) = \text{cycle time}; 1 / (500 \text{ MHz}) = 2 \text{ ns}$





MIPS and MFLOPS

- MIPS and MFLOPS are common metrics for describing the capabilities of the machines
 - But, are they suitable for comparing performance of different machines all the time?
- MIPS: million instructions per second
$$= (\text{number of } \mathbf{\text{instructions}} \text{ executed in program}) / (\text{execution time in seconds} * 10^6)$$
- MFLOPS: million floating point operations per second
$$= (\text{number of } \mathbf{\text{floating point operations}} \text{ executed in program}) / (\text{execution time in seconds} * 10^6)$$





A Benchmarking Example

- Pentium III 2.5Ghz system, Microsoft C++ compiler
 - Compile the program, execute, count instructions
 - Measured at **2,100 MIPS**
- What does this tell you about performance?
- Compile again, this time with optimization **ON!**
 - Compilation takes a lot longer, execute, count instructions
 - Measured performance at **1,600 MIPS**
- Intuitively, the optimized version should be faster than the original version (and in fact, the optimized version takes less time)
- But, the MIPS of the former is less than that of the latter. What happened?





A Benchmarking Example (Cont'd)

of Inst_A /
ExecTime_A

of Inst_B /
ExecTime_B

- There are **fewer instructions** executed in the optimized program!
 - And the optimized program takes less time
- MIPS rating depends on **compiler**
 - Quality of generated code
 - **Optimized for instruction execution time, not MIPS rating**
 - Compilers are always benchmarked with the machine
- How could you “*cheat*” to get a high MIPS rating?





Another Benchmarking Example

- Power Macintosh, 1Ghz, PowerPC G4
 - Compile same program, “optimized”
 - Execute, assuming no obvious cheating
 - Experiment produces 1,500 MIPS rating
 - Is this faster than the Pentium III?

=> There's no easy way to tell from this information!

- Why?
 - The **unit of work** has changed.
 - **Pentium Instruction != PowerPC instruction**
- **Hard to compare MIPS across architectures**
 - MIPS is of little use for comparing architectures
 - Resort to execution time





Another Benchmarking Example (Cont'd)

- **Unit of work**
 - **The real target** when we try to benchmark a program/system
 - E.g., Instructions, Floating Point Operations, Window updates, etc.
- The benchmarking result can be considered as the time taken for executing **the works**, which are affected by different factors:
 - Instructions: compiler, architecture
 - Floating Point operations: compiler, algorithm
 - Window updates: algorithm
- That is, the benchmarking result is the combined effect of the HW and SW environment
 - **Depends** on compiler, architecture, algorithm, implementation, execution environment, etc.





Example: Unit of Work

- Floating Point Operations
- Window Updates
- Frames/Polygons (rendering)
- Megabytes (communication)
- Limitations of each of these?
- How can you cheat/reduce each of these?
 - Time to finish a fixed amount of work
 - Ratio of the finished amount of work per time unit

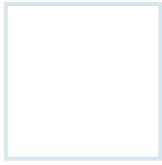
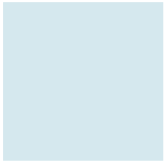




Summary

- Architecture involves a tension of programmability, compilers, implementability, and technology
 - Optimize the design given these ever changing constraints
- Many possible measures of **work** / **performance metrics**
- Choosing is rife with potential errors
- Because it includes everything, **execution time** is the safest choice
- Still need to analyze the other influences carefully before you can draw any conclusions about the causes






TIME MEASUREMENT FOR PERFORMANCE ANALYSIS

Instrumentation-based Time Measurement



- Code instrumentation
 - Inject the *time functions* into the target application
 - It can be done manually or automatically in the source or binary program
- Example
 - Insert *time* at the prolog and epilog of each function to get the elapsed times of the functions

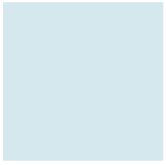




Watch Out for Its Overhead

- Code instrumentation overhead
 - The **red blocks** refer to the **overhead** incurred by the **time functions**
 - The overhead is often ignored; but, you should pay attention to it
 - Here, we assume it is negligible compared with the **targets** (e.g., the functions to be analyzed)
- Instrumentation always:
 - **introduces overhead** (low or high)
 - **alters the program execution**





Measurement of *Time* Overhead

- It is crucial to understand the **cost** of the time measurement
- You may use the following simple method to get the **overhead** of the time function you used
 - Example below shows the averaged time required by each **time()** function invocation
- Note: **one million** is a magic number

```
start = time();  
for i= 1 ~ 1,000,000  
    time();  
end = time();
```

overhead = (end-start)/1,000,000

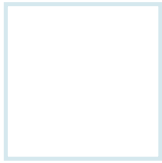
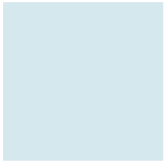




Note

- The following methods are simple and good for the **sequential** programs
- As for **parallel programs**, you should consider the concurrency issues
 - E.g., Lock and Wait
 - In this case, performance profiling tools should be used to characterize the runtime activities of multiple threads





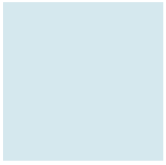
STANDARD TIME MEASUREMENT



Be Careful

- The length of code for time measurement
- Unit & Resolution
 - Minute, seconds, microseconds
 - Length of code vs. Time Resolution
- Variables to hold the *time* values
 - Each time function has its own rules and way to *represent* the time
 - Types/Formats of the variables are important
- *Clock source* of the time function
 - System-wide, per-process, per-thread
- Code portability
 - Is the time measurement code portable across HW platforms?
 - E.g., Windows system has its time functions





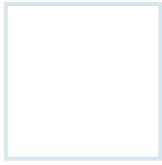
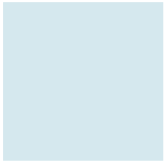
time()

- Get time in seconds
 - Return the current calendar time or -1 if there is an error
 - If the argument time is given, then the current time is stored in time
 - Only measure the time up to seconds

```
#include <ctime>
```

```
time_t time( time_t *time );
```





clock()

- Determine processor time (clocks)
 - Return the processes CPU time (time since the program started)
 - Return -1 if that information is not available
- Conversion to seconds by division by `CLOCKS_PER_SEC`
 - Note: if your compiler is POSIX compliant, then `CLOCKS_PER_SEC` is always defined as 1,000,000
 - On a 32-bit system where `CLOCKS_PER_SEC` equals 1000000 this function will **return the same value approximately every 72 minutes**
 - For improved accuracy, since glibc 2.18, it is implemented on top of `clock_gettime(2)` (using the `CLOCK_PROCESS_CPUTIME_ID` clock)

```
#include <ctime>
```

```
clock_t clock( void );
```





Example: Time Measurement using clock()

```
#include <ctime>
```

```
// Time stamp before the computations
```

```
clock_t start = clock();
```

```
... /* Computations to be measured */
```

```
// Time stamp after the computations
```

```
clock_t end = clock();
```

```
double cpu_time = static_cast<double>( end - start ) /
```

```
CLOCKS_PER_SEC;
```

Report the elapsed time for the code section enclosed by the two **clock** functions





clock_gettime()

- High Precision Event Timer
- The functions **clock_gettime()** and **clock_settime()** retrieve and set the time of the specified clock
- The function **clock_getres()** finds the resolution (precision) of the specified clock *clk_id* (should also consider the incurred overhead)

```
#include <time.h>
```

```
int clock_gettime(clockid_t clk_id, struct timespec *tp);
```

```
struct timespec {
```

```
    time_t  tv_sec;      /* seconds */
```

```
    long    tv_nsec;     /* nanoseconds */
```

```
};
```

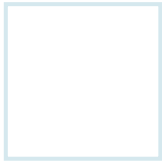
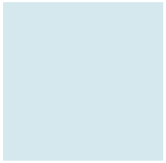




Supported Clocks for `clock_gettime()`

- The **clk_id** argument is the identifier of the particular clock on which to act
- Sufficiently recent versions of GNU libc and the Linux kernel support the following clocks:
 - **CLOCK_REALTIME**
System-wide wall-clock clock. Set this clock requiring appropriate privileges
 - **CLOCK_MONOTONIC**
Clock that cannot be set and represents monotonic time since some unspecified starting point
Good for total elapsed time, including I/O & block overhead
 - **CLOCK_PROCESS_CPUTIME_ID**
High-resolution per-process timer from the CPU (not for block/sleep code)
 - **CLOCK_THREAD_CPUTIME_ID**
Thread-specific CPU-time clock





NATIVE LINUX TIME MEASUREMENT





gettimeofday()

Used for measuring wall clock time

```
#include <sys/time.h>
#include <sys/types.h>
struct timeval tp;
double sec, usec, start, end;
// Time stamp before the computations
gettimeofday( &tp, NULL );
sec = static_cast<double>( tp.tv_sec );
usec = static_cast<double>( tp.tv_usec )/1E6;
start = sec + usec;
// Computations to be measured
...
...
// Time stamp after the computations
gettimeofday( &tp, NULL );
sec = static_cast<double>( tp.tv_sec );
usec = static_cast<double>( tp.tv_usec )/1E6;
end = sec + usec;
// Time calculation (in seconds)
double time = end - start;
```





TIME MEASUREMENT FOR X86 PLATFORMS





Read Time Stamp Counter (RDTSC)

- Time Stamp Counter (TSC)
 - Record CPU clock cycles
 - Monotonically increase for each CPU clock
 - Accessed by the Read Time Stamp Counter (**RDTSC**) instruction introduced in Pentium processors
- Simple and neat to use the Time Stamp Counter for time measurement
 - by calling the inline function below to get the current timestamp (in cycles)
 - Need to translate the cycles into the actual time

/* assembly code to read the TSC */

static inline uint64_t RDTSC()

```
{
    unsigned int hi, lo;
    __asm__ volatile("rdtsc" : "=a" (lo), "=d" (hi));
    return ((uint64_t)hi << 32) | lo;
}
```

Courtesy of <http://www.intel.com/content/dam/www/public/us/en/documents/white-papers/ia-32-ia-64-benchmark-code-execution-paper.pdf>
<http://stackoverflow.com/questions/3388134/rdtsc-accuracy-across-cpu-cores>

Also read: <http://oliveryang.net/2015/09/pitfalls-of-TSC-usage/>

November 22, 2022





Why RDTSC?

- Compared with **clock_gettime()**
 - Higher resolution
 - We can get the resolution through the API **clock_getres()**
 - On the Dell XPS 1530 with Intel core2duo T7500 CPU running Ubuntu 10.04, it has a resolution of 1 nanosecond
 - On the other hand, RDTSC instruction can have resolution of up to a CPU clock time
 - On the 2.2 GHz CPU that means resolution is 0.45 nanoseconds
 - Low cost
 - Measure the time taken for 1 million calls to both HPET and RDTSC
 - HPET : 1 s 482 ms 188 us 38 ns
 - RDTSC: 0 s 103 ms 311 us 752 ns
 - RDTSC is 14x faster than HPET

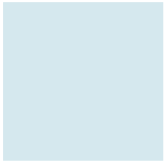




Be Careful for the Following Situations!!!

- The per-CPU TSC value may not be the same across CPU cores
 - Multiple cores having different TSC values
 - One should bind the process to same CPU
 - Invariant TSC support for ensuring consistent across multiple cores; but, one should pay attention to multi-CPU
- CPU frequency scaling for power saving
 - Fixed CPU power governor policy, e.g., high performance
- Hibernation of system will reset TSC value
 - Disable the hibernate; or, check the TSC values
- Impact on portability due to varying implementation of CPUs
 - Fixed Intel CPUs with same settings
- Out-of-order execution of code
 - Use the instruction, e.g., CPUID, for serializing instructions
 - Please refer to [the document](#)





QUESTIONS?

